

Prediction_Assignment_Writeup

2015-09-27

```
library(caret)
library(rpart)
library(rattle)
library(randomForest)
library(knitr)
```

Project

Project Introduction

Background

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. These type of devices are part of the quantified self movement - a group of enthusiasts who take measurements about themselves regularly to improve their health, to find patterns in their behavior, or because they are tech geeks. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. In this project, your goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways. More information is available from the website here: <http://groupware.les.inf.puc-rio.br/har> (see the section on the Weight Lifting Exercise Dataset).

Data

The training data for this project are available here:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>

The test data are available here:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>

The data for this project come from this source: <http://groupware.les.inf.puc-rio.br/har>. If you use the document you create for this class for any purpose please cite them as they have been very generous in allowing their data to be used for this kind of assignment.

Goal

The goal of your project is to predict the manner in which they did the exercise. This is the “classe” variable in the training set. You may use any of the other variables to

predict with. You should create a report describing how you built your model, how you used cross validation, what you think the expected out of sample error is, and why you made the choices you did. You will also use your prediction model to predict 20 different test cases.

Getting And Loading The Data

```
set.seed(12345)

trainUrl <- "http://d396qusza40orc.cloudfront.net/predmachlearn/pml-
training.csv"
testUrl <- "http://d396qusza40orc.cloudfront.net/predmachlearn/pml-
testing.csv"

training <- read.csv(url(trainUrl), na.strings=c("NA", "#DIV/0!", ""))
testing <- read.csv(url(testUrl), na.strings=c("NA", "#DIV/0!", ""))

colnames_train <- colnames(training)
```

Cleaning The Data

The first step was to clean the data from all kind of missing values and columns that may be irrelevant to prediction (i.e - near zero variance columns) ###Removing Columns With NAs

```
# Count the number of non-NAs in each col.
nonNAs <- function(x) {
  as.vector(apply(x, 2, function(x) length(which(!is.na(x)))))
}

# Build vector of missing data or NA columns to drop.
colcnts <- nonNAs(training)
drops <- c()
for (cnt in 1:length(colcnts)) {
  if (colcnts[cnt] < nrow(training)) {
    drops <- c(drops, colnames_train[cnt])
  }
}
```

Removing Irrelevant Columns

These columns contains irrelevant information for the prediction algorithm.

```
# Drop NA data and the first 7 columns as they're unnecessary for
predicting.
training <- training[,!(names(training) %in% drops)]
training <- training[,8:length(colnames(training))]
```

```
testing <- testing[,!(names(testing) %in% drops)]
testing <- testing[,8:length(colnames(testing))]
```

Show Remaining Columns Training vs. Testing

Show remaining columns training.

```
colnames(training)
```

```
## [1] "roll_belt"          "pitch_belt"         "yaw_belt"
## [4] "total_accel_belt"   "gyros_belt_x"       "gyros_belt_y"
## [7] "gyros_belt_z"       "accel_belt_x"       "accel_belt_y"
## [10] "accel_belt_z"       "magnet_belt_x"      "magnet_belt_y"
## [13] "magnet_belt_z"      "roll_arm"           "pitch_arm"
## [16] "yaw_arm"            "total_accel_arm"    "gyros_arm_x"
## [19] "gyros_arm_y"        "gyros_arm_z"        "accel_arm_x"
## [22] "accel_arm_y"        "accel_arm_z"        "magnet_arm_x"
## [25] "magnet_arm_y"       "magnet_arm_z"       "roll_dumbbell"
## [28] "pitch_dumbbell"     "yaw_dumbbell"
"total_accel_dumbbell"
## [31] "gyros_dumbbell_x"   "gyros_dumbbell_y"
"gyros_dumbbell_z"
## [34] "accel_dumbbell_x"   "accel_dumbbell_y"
"accel_dumbbell_z"
## [37] "magnet_dumbbell_x"  "magnet_dumbbell_y"
"magnet_dumbbell_z"
## [40] "roll_forearm"       "pitch_forearm"      "yaw_forearm"
## [43] "total_accel_forearm" "gyros_forearm_x"    "gyros_forearm_y"
## [46] "gyros_forearm_z"    "accel_forearm_x"    "accel_forearm_y"
## [49] "accel_forearm_z"    "magnet_forearm_x"
"magnet_forearm_y"
## [52] "magnet_forearm_z"   "classe"
```

Show remaining columns testing

```
colnames(testing)
```

```
## [1] "roll_belt"          "pitch_belt"         "yaw_belt"
## [4] "total_accel_belt"   "gyros_belt_x"       "gyros_belt_y"
## [7] "gyros_belt_z"       "accel_belt_x"       "accel_belt_y"
## [10] "accel_belt_z"       "magnet_belt_x"      "magnet_belt_y"
## [13] "magnet_belt_z"      "roll_arm"           "pitch_arm"
## [16] "yaw_arm"            "total_accel_arm"    "gyros_arm_x"
## [19] "gyros_arm_y"        "gyros_arm_z"        "accel_arm_x"
## [22] "accel_arm_y"        "accel_arm_z"        "magnet_arm_x"
## [25] "magnet_arm_y"       "magnet_arm_z"       "roll_dumbbell"
## [28] "pitch_dumbbell"     "yaw_dumbbell"
"total_accel_dumbbell"
## [31] "gyros_dumbbell_x"   "gyros_dumbbell_y"
"gyros_dumbbell_z"
## [34] "accel_dumbbell_x"   "accel_dumbbell_y"
"accel_dumbbell_z"
## [37] "magnet_dumbbell_x"  "magnet_dumbbell_y"
```

```

"magnet_dumbbell_z"
## [40] "roll_forearm"      "pitch_forearm"      "yaw_forearm"
## [43] "total_accel_forearm" "gyros_forearm_x"    "gyros_forearm_y"
## [46] "gyros_forearm_z"    "accel_forearm_x"    "accel_forearm_y"
## [49] "accel_forearm_z"    "magnet_forearm_x"
"magnet_forearm_y"
## [52] "magnet_forearm_z"    "problem_id"

```

Removing Columns With Near Zero Variance

```
print(nearZeroVar(training, saveMetrics=TRUE))
```

```

##              freqRatio percentUnique zeroVar  nzv
## roll_belt      1.101904      6.7781062  FALSE FALSE
## pitch_belt     1.036082      9.3772296  FALSE FALSE
## yaw_belt       1.058480      9.9734991  FALSE FALSE
## total_accel_belt 1.063160      0.1477933  FALSE FALSE
## gyros_belt_x    1.058651      0.7134849  FALSE FALSE
## gyros_belt_y    1.144000      0.3516461  FALSE FALSE
## gyros_belt_z    1.066214      0.8612782  FALSE FALSE
## accel_belt_x    1.055412      0.8357966  FALSE FALSE
## accel_belt_y    1.113725      0.7287738  FALSE FALSE
## accel_belt_z    1.078767      1.5237998  FALSE FALSE
## magnet_belt_x   1.090141      1.6664968  FALSE FALSE
## magnet_belt_y   1.099688      1.5187035  FALSE FALSE
## magnet_belt_z   1.006369      2.3290184  FALSE FALSE
## roll_arm       52.338462     13.5256345  FALSE FALSE
## pitch_arm      87.256410     15.7323412  FALSE FALSE
## yaw_arm        33.029126     14.6570176  FALSE FALSE
## total_accel_arm 1.024526      0.3363572  FALSE FALSE
## gyros_arm_x     1.015504      3.2769341  FALSE FALSE
## gyros_arm_y     1.454369      1.9162165  FALSE FALSE
## gyros_arm_z     1.110687      1.2638875  FALSE FALSE
## accel_arm_x     1.017341      3.9598410  FALSE FALSE
## accel_arm_y     1.140187      2.7367241  FALSE FALSE
## accel_arm_z     1.128000      4.0362858  FALSE FALSE
## magnet_arm_x    1.000000      6.8239731  FALSE FALSE
## magnet_arm_y    1.056818      4.4439914  FALSE FALSE
## magnet_arm_z    1.036364      6.4468454  FALSE FALSE
## roll_dumbbell   1.022388     84.2065029  FALSE FALSE
## pitch_dumbbell  2.277372     81.7449801  FALSE FALSE
## yaw_dumbbell    1.132231     83.4828254  FALSE FALSE
## total_accel_dumbbell 1.072634      0.2191418  FALSE FALSE
## gyros_dumbbell_x 1.003268      1.2282132  FALSE FALSE
## gyros_dumbbell_y 1.264957      1.4167771  FALSE FALSE
## gyros_dumbbell_z 1.060100      1.0498420  FALSE FALSE
## accel_dumbbell_x 1.018018      2.1659362  FALSE FALSE
## accel_dumbbell_y 1.053061      2.3748853  FALSE FALSE
## accel_dumbbell_z 1.133333      2.0894914  FALSE FALSE
## magnet_dumbbell_x 1.098266      5.7486495  FALSE FALSE

```

## magnet_dumbbell_y	1.197740	4.3012945	FALSE	FALSE
## magnet_dumbbell_z	1.020833	3.4451126	FALSE	FALSE
## roll_forearm	11.589286	11.0895933	FALSE	FALSE
## pitch_forearm	65.983051	14.8557741	FALSE	FALSE
## yaw_forearm	15.322835	10.1467740	FALSE	FALSE
## total_accel_forearm	1.128928	0.3567424	FALSE	FALSE
## gyros_forearm_x	1.059273	1.5187035	FALSE	FALSE
## gyros_forearm_y	1.036554	3.7763735	FALSE	FALSE
## gyros_forearm_z	1.122917	1.5645704	FALSE	FALSE
## accel_forearm_x	1.126437	4.0464784	FALSE	FALSE
## accel_forearm_y	1.059406	5.1116094	FALSE	FALSE
## accel_forearm_z	1.006250	2.9558659	FALSE	FALSE
## magnet_forearm_x	1.012346	7.7667924	FALSE	FALSE
## magnet_forearm_y	1.246914	9.5403119	FALSE	FALSE
## magnet_forearm_z	1.000000	8.5771073	FALSE	FALSE
## classe	1.469581	0.0254816	FALSE	FALSE

No headers with nzc were found.


Partitioning The Training Set

```
set.seed(666)
inTrain <- createDataPartition(y=training$classe, p=0.6, list=FALSE)
df_training <- training[inTrain,]
df_testing <- training[-inTrain,]
```

Prediction With Decision Tree

```
modFitA <- rpart(classe ~ ., data=df_training, method="class")
fancyRpartPlot(modFitA)
```

```
## Warning: labs do not fit even at cex 0.15, there may be some
overplotting
```



Rattle 2015-Sep-27 16:48:24 okatz

```
predictionsA <- predict(modFitA, df_testing, type = "class")
confusionMatrix(predictionsA, df_testing$classe)
```

Confusion Matrix and Statistics

```
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1970  344   44  131   99
##           B   85  890  140   49  133
##           C   45  106 1001  178  132
##           D  105  100   99  856  142
##           E   27   78   84   72  936
```

##

Overall Statistics

##

```
##           Accuracy : 0.7205
##           95% CI : (0.7104, 0.7304)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
```

##

```
##           Kappa : 0.6446
##           Mcnemar's Test P-Value : < 2.2e-16
```

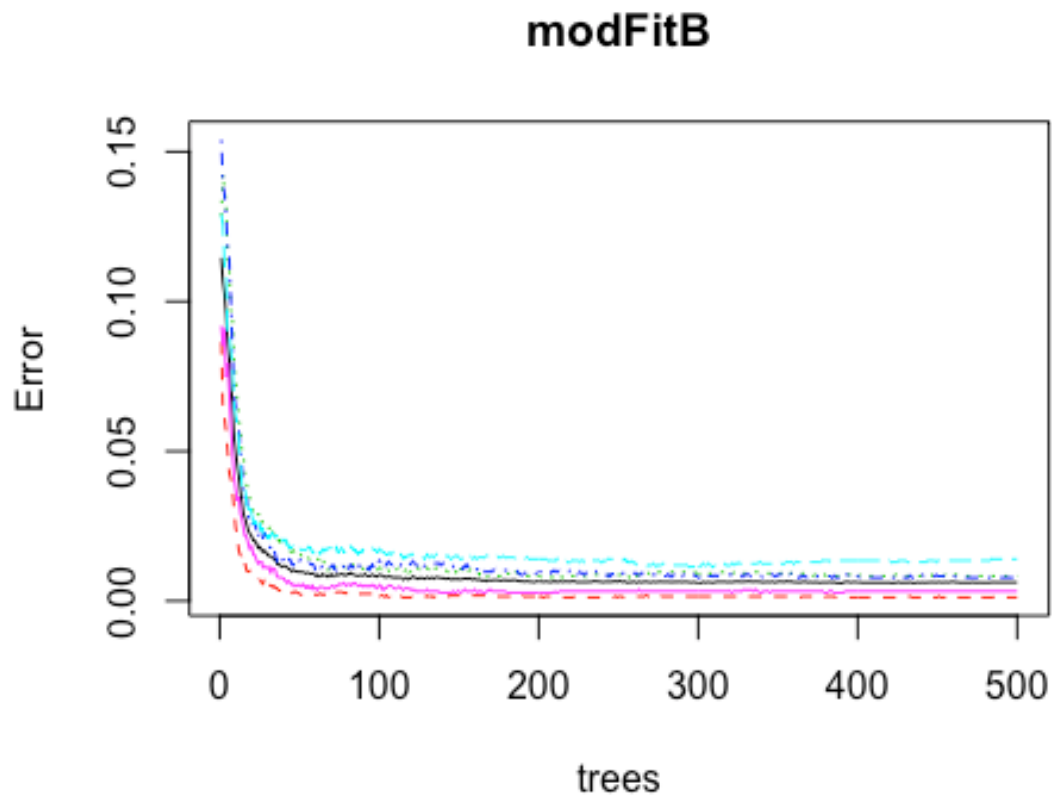
##

```
## Statistics by Class:
##
##               Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.8826   0.5863   0.7317   0.6656   0.6491
## Specificity      0.8899   0.9357   0.9288   0.9320   0.9592
## Pos Pred Value   0.7612   0.6862   0.6847   0.6575   0.7820
## Neg Pred Value    0.9502   0.9041   0.9425   0.9343   0.9239
## Prevalence       0.2845   0.1935   0.1744   0.1639   0.1838
## Detection Rate    0.2511   0.1134   0.1276   0.1091   0.1193
## Detection Prevalence 0.3298 0.1653 0.1863 0.1659 0.1526
## Balanced Accuracy 0.8863 0.7610 0.8303 0.7988 0.8042
```

The accuracy is not good enough 72% therefore I also tried Random forest algo to see if we can find better prediction.

Prediction With Random Forest

```
set.seed(666)
modFitB <- randomForest(classe ~. , data=df_training)
plot(modFitB)
```



```
predictionsB <- predict(modFitB, df_testing, type = "class")
confusionMatrix(predictionsB, df_testing$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 2229   12    0    0    0
##           B    3 1500   20    0    0
##           C    0    6 1345   14    0
##           D    0    0    3 1272    5
##           E    0    0    0    0 1437
##
## Overall Statistics
##
##           Accuracy : 0.992
##           95% CI : (0.9897, 0.9938)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9898
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9987   0.9881   0.9832   0.9891   0.9965
## Specificity      0.9979   0.9964   0.9969   0.9988   1.0000
## Pos Pred Value   0.9946   0.9849   0.9853   0.9938   1.0000
## Neg Pred Value   0.9995   0.9972   0.9965   0.9979   0.9992
## Prevalence       0.2845   0.1935   0.1744   0.1639   0.1838
## Detection Rate   0.2841   0.1912   0.1714   0.1621   0.1832
## Detection Prevalence 0.2856   0.1941   0.1740   0.1631   0.1832
## Balanced Accuracy 0.9983   0.9923   0.9900   0.9939   0.9983
```

Random forest yield better results with 99% accuracy!

Assignment Submission & Result Prediction

Random Forests gave an Accuracy on the training dataset of 99.2%, which was more accurate than what I got from the Decision Tree with 72.05%. The expected out-of-sample error is $100 - 99.2 = 0.8\%$.

```
predictionsTest <- predict(modFitB, testing, type = "class")

predictionsTest

##  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20
##  B  A  B  A  A  E  D  B  A  A  B  C  B  A  E  E  A  B  B  B
## Levels: A B C D E

pml_write_files = function(x){
  n = length(x)
  for(i in 1:n){
```



```
        filename = paste0("problem_id_",i, ".txt")
write.table(x[i],file=filename,quote=FALSE,row.names=FALSE,col.names=FALSE)
    }
}

pml_write_files(predictionsTest)
```